

Submitted By

- Ruthvic Punyamurtula (Problems - 3,5)
- Sai Charan Kothapalli (problems - 1,2,4)

Resources

- [youtube demo - 1,2,4](#)
- [youtube demo - 3,5](#)
- [Source Code](#)

Introduction

This is lab assignment 3 of CS5590 - python/Deep Learning class. This lab specifically consists of several deep learning tasks based on regression, CNN and LSTM.

Objective

In this assignment, we used Kaggle datasets & implemented

- Linear Regression
- Logistics Regression
- Image Classification with CNN
- Text Classification with CNN
- Text classification with LSTM
- Compare Text Classification with CNN & LSTM

Approaches/Methods

1.Linear Regression

```
using tensorflow backend.
```

	crim	zn	...	lstat	medv
count	506.000000	506.000000	...	506.000000	506.000000
mean	3.613524	11.363636	...	12.653063	22.532806
std	8.601545	23.322453	...	7.141062	9.197104
min	0.006320	0.000000	...	1.730000	5.000000
25%	0.082045	0.000000	...	6.950000	17.025000
50%	0.256510	0.000000	...	11.360000	21.200000
75%	3.677082	12.500000	...	16.955000	25.000000
max	88.976200	100.000000	...	37.970000	50.000000

```
[8 rows x 14 columns]
Training shape: (354, 13)
ddd (354,)
Training samples: 354
Validation samples: 152
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 100)	1400
dense_2 (Dense)	(None, 50)	5050
dense_3 (Dense)	(None, 1)	51

```
Total params: 6,501
Trainable params: 6,501
Non-trainable params: 0
```

None

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 100)	1400
dense_2 (Dense)	(None, 50)	5050
dense_3 (Dense)	(None, 1)	51

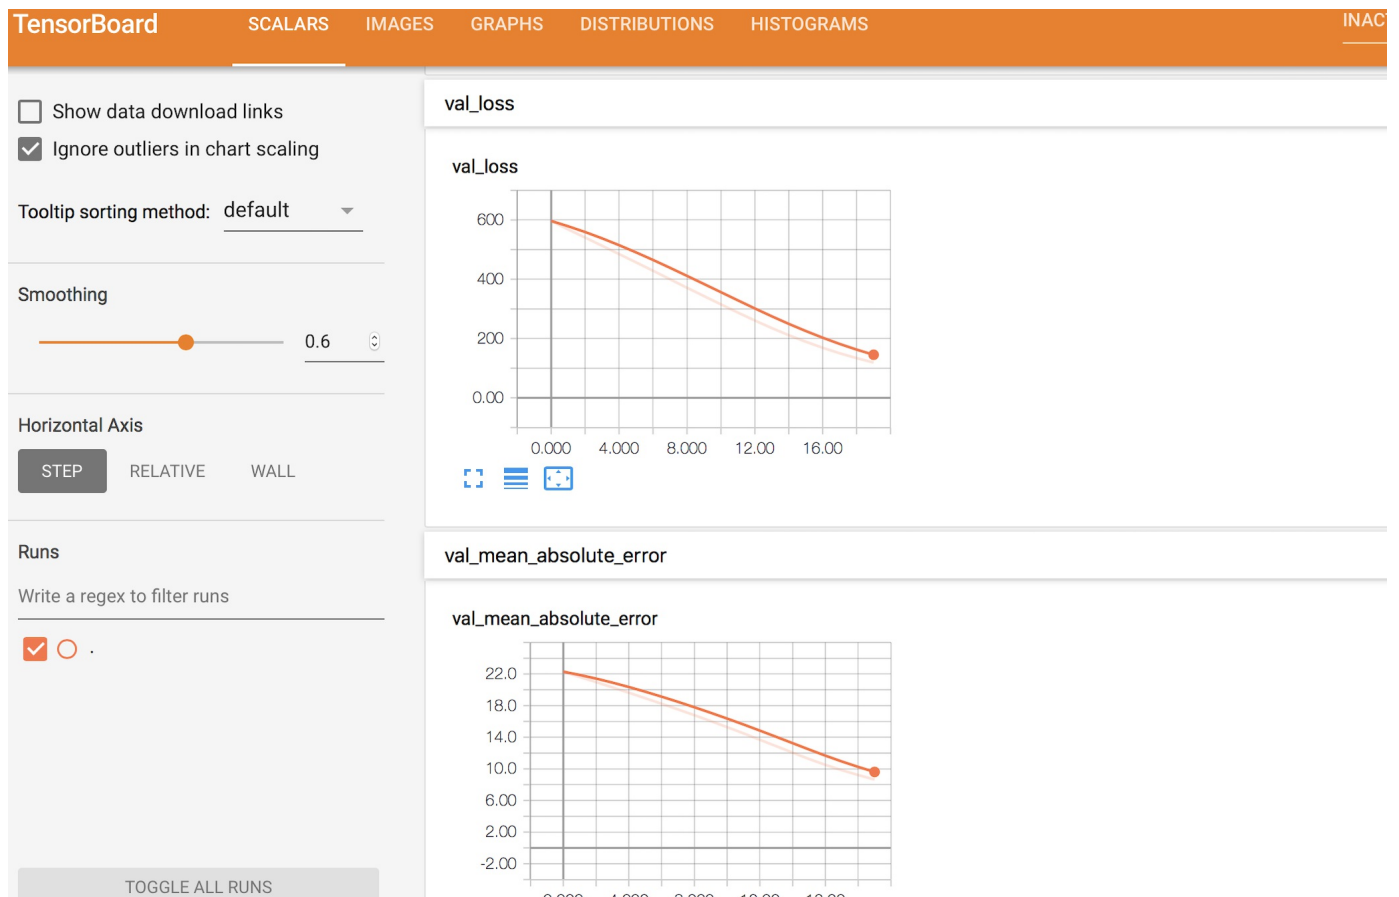
```
Total params: 6,501
Trainable params: 6,501
Non-trainable params: 0
```

Train on 354 samples. validate on 152 samples

```
Epoch 1/20
- 0s - loss: 553.8860 - mean_absolute_error: 21.8283 - val_loss: 596.6977 - val_mean_absolute_error: 22.2960
Epoch 2/20
- 0s - loss: 521.4080 - mean_absolute_error: 21.0662 - val_loss: 567.9603 - val_mean_absolute_error: 21.6311
Epoch 3/20
- 0s - loss: 493.8241 - mean_absolute_error: 20.3940 - val_loss: 540.6460 - val_mean_absolute_error: 20.9797
Epoch 4/20
- 0s - loss: 466.9822 - mean_absolute_error: 19.7284 - val_loss: 512.8830 - val_mean_absolute_error: 20.3116
Epoch 5/20
- 0s - loss: 440.0941 - mean_absolute_error: 19.0482 - val_loss: 484.9061 - val_mean_absolute_error: 19.6317
Epoch 6/20
- 0s - loss: 413.1259 - mean_absolute_error: 18.3552 - val_loss: 456.9809 - val_mean_absolute_error: 18.9385
Epoch 7/20
- 0s - loss: 386.1872 - mean_absolute_error: 17.6299 - val_loss: 428.7061 - val_mean_absolute_error: 18.2341
Epoch 8/20
- 0s - loss: 359.1472 - mean_absolute_error: 16.8844 - val_loss: 400.0366 - val_mean_absolute_error: 17.5226
Epoch 9/20
- 0s - loss: 332.0476 - mean_absolute_error: 16.1182 - val_loss: 371.2846 - val_mean_absolute_error: 16.7870
Epoch 10/20
- 0s - loss: 305.3450 - mean_absolute_error: 15.3358 - val_loss: 343.3200 - val_mean_absolute_error: 16.0444
Epoch 11/20
- 0s - loss: 279.2301 - mean_absolute_error: 14.5623 - val_loss: 315.3151 - val_mean_absolute_error: 15.2790
Epoch 12/20
- 0s - loss: 253.5218 - mean_absolute_error: 13.7447 - val_loss: 287.6593 - val_mean_absolute_error: 14.4908
Epoch 13/20
- 0s - loss: 228.4915 - mean_absolute_error: 12.9200 - val_loss: 260.7425 - val_mean_absolute_error: 13.6798
Epoch 14/20
- 0s - loss: 204.6113 - mean_absolute_error: 12.0920 - val_loss: 235.2078 - val_mean_absolute_error: 12.8644
Epoch 15/20
- 0s - loss: 182.4678 - mean_absolute_error: 11.2679 - val_loss: 211.3075 - val_mean_absolute_error: 12.0524
Epoch 16/20
- 0s - loss: 162.2734 - mean_absolute_error: 10.4894 - val_loss: 189.3818 - val_mean_absolute_error: 11.2484
Epoch 17/20
- 0s - loss: 143.8153 - mean_absolute_error: 9.7355 - val_loss: 168.8125 - val_mean_absolute_error: 10.5032
Epoch 18/20
- 0s - loss: 127.3597 - mean_absolute_error: 9.0520 - val_loss: 150.3502 - val_mean_absolute_error: 9.8346
Epoch 19/20
- 0s - loss: 112.8586 - mean_absolute_error: 8.4007 - val_loss: 133.9125 - val_mean_absolute_error: 9.2193
Epoch 20/20
- 0s - loss: 100.4432 - mean_absolute_error: 7.8198 - val_loss: 119.1148 - val_mean_absolute_error: 8.6499
Train MAE: 7.437 , Train Loss: 92.4845
Val MAE: 8.6499 , Val Loss: 119.1148
```

Process finished with exit code 0

Tensorboard



2. Logistic Regression

Using TensorFlow backend.

	Time	V1	...	Amount	Class
count	284807.000000	2.848070e+05	...	284807.000000	284807.000000
mean	94813.859575	1.165980e-15	...	88.349619	0.001727
std	47488.145955	1.958696e+00	...	250.120109	0.041527
min	0.000000	-5.640751e+01	...	0.000000	0.000000
25%	54201.500000	-9.203734e-01	...	5.600000	0.000000
50%	84692.000000	1.810880e-02	...	22.000000	0.000000
75%	139320.500000	1.315642e+00	...	77.165000	0.000000
max	172792.000000	2.454930e+00	...	25691.160000	1.000000

[8 rows x 31 columns]

Training shape: (199364, 20)

ddd (199364,)

Training samples: 199364

Validation samples: 85443

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 100)	2100
dropout_1 (Dropout)	(None, 100)	0
dense_2 (Dense)	(None, 50)	5050
dense_3 (Dense)	(None, 20)	1020
dense_4 (Dense)	(None, 1)	21

Total params: 8,191

Trainable params: 8,191

Non-trainable params: 0

None

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 100)	2100
dropout_1 (Dropout)	(None, 100)	0
dense_2 (Dense)	(None, 50)	5050
dense_3 (Dense)	(None, 20)	1020

Using TensorFlow backend.

	Time	V1	...	Amount	Class
count	284807.000000	2.848070e+05	...	284807.000000	284807.000000
mean	94813.859575	1.165980e-15	...	88.349619	0.001727
std	47488.145955	1.958696e+00	...	250.120109	0.041527
min	0.000000	-5.640751e+01	...	0.000000	0.000000
25%	54201.500000	-9.203734e-01	...	5.600000	0.000000
50%	84692.000000	1.810880e-02	...	22.000000	0.000000
75%	139320.500000	1.315642e+00	...	77.165000	0.000000
max	172792.000000	2.454930e+00	...	25691.160000	1.000000

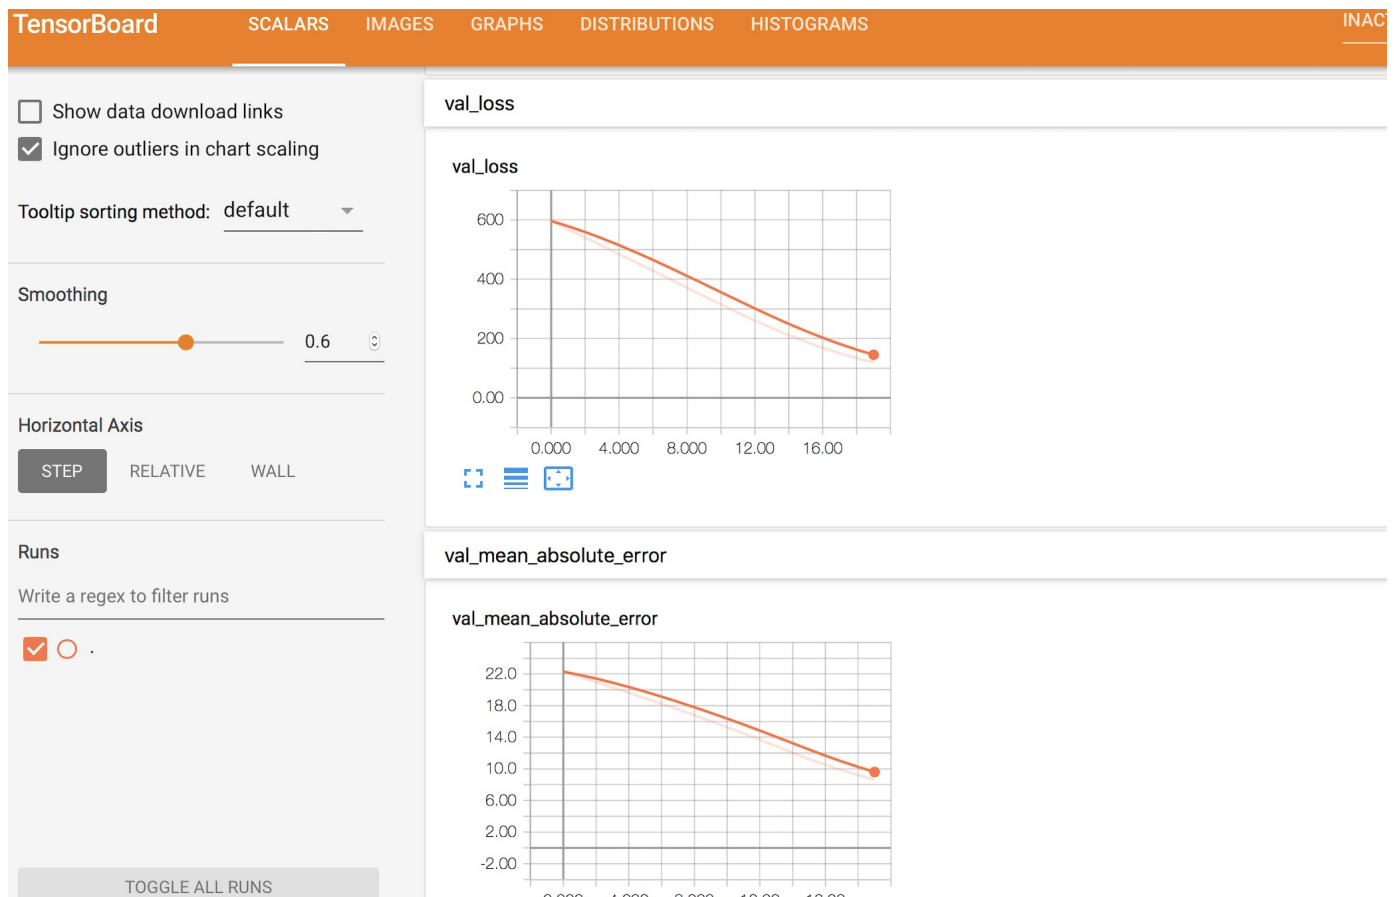
[8 rows x 31 columns]
Training shape: (199364, 20)
ddd (199364,)
Training samples: 199364
Validation samples: 85443

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 100)	2100
dropout_1 (Dropout)	(None, 100)	0
dense_2 (Dense)	(None, 50)	5050
dense_3 (Dense)	(None, 20)	1020
dense_4 (Dense)	(None, 1)	21

Total params: 8,191
Trainable params: 8,191
Non-trainable params: 0

None		
Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 100)	2100
dropout_1 (Dropout)	(None, 100)	0
dense_2 (Dense)	(None, 50)	5050
dense_3 (Dense)	(None, 20)	1020

Tensorboard



3. Image Classification with CNN

The dataset used is - <https://www.kaggle.com/slothkong/10-monkey-species>

The initial data had 10 classes/species of monkeys. Due to laptop performance issues, we have reduced the dataset to 3 species of monkeys and loaded the data into /content of google colab - with runtime as GPU.

The code snippet describing the data size is given below


```
# Training generator
train_datagen = ImageDataGenerator(rotation_range = 30
                                   ,rescale=1. / 255,
                                   shear_range=0.2,
                                   zoom_range=0.2,
                                   horizontal_flip=True)
train_generator = train_datagen.flow_from_directory(train_dir,
                                                    target_size=(height,width),
                                                    batch_size=batch_size,
                                                    seed=seed,
                                                    class_mode='categorical')
```

Found 326 images belonging to 3 classes.

```
7]: # Test generator
test_datagen = ImageDataGenerator(rescale=1./255)
test_generator = test_datagen.flow_from_directory(test_dir,
                                                  target_size=(height,width),
                                                  batch_size=batch_size,
                                                  seed=seed,
                                                  class_mode='categorical')
```

Found 81 images belonging to 3 classes.

The model used for classification is


Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 148, 148, 32)	896
flatten_1 (Flatten)	(None, 700928)	0
dense_1 (Dense)	(None, 512)	358875648
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 256)	131328
dense_3 (Dense)	(None, 3)	771
Total params: 359,008,643		
Trainable params: 359,008,643		
Non-trainable params: 0		

Because we had small data & training on 20 epochs the accuracy

obtained is low (about 35%)

```
Epoch 15/20
8/8 [=====] - 25s 3s/step - loss: 10.4578 - acc: 0.3512 - val_loss: 10.5464 - val_acc: 0.3457
Epoch 16/20
8/8 [=====] - 24s 3s/step - loss: 10.5652 - acc: 0.3445 - val_loss: 10.5464 - val_acc: 0.3457
Epoch 17/20
8/8 [=====] - 23s 3s/step - loss: 10.4996 - acc: 0.3486 - val_loss: 10.5464 - val_acc: 0.3457
Epoch 18/20
8/8 [=====] - 25s 3s/step - loss: 10.6935 - acc: 0.3366 - val_loss: 10.5464 - val_acc: 0.3457
Epoch 19/20
8/8 [=====] - 25s 3s/step - loss: 10.6505 - acc: 0.3392 - val_loss: 10.5464 - val_acc: 0.3457
Epoch 20/20
8/8 [=====] - 23s 3s/step - loss: 10.7096 - acc: 0.3356 - val_loss: 10.5464 - val_acc: 0.3457
```

The plots showing accuracy & loss for train & validation data is shown below



4. Text Classification with CNN

```
8320/8529 [=====>.] - ETA: 1s - loss: 0.0155 - acc: 0.9996
8400/8529 [=====>.] - ETA: 1s - loss: 0.0155 - acc: 0.9996
8480/8529 [=====>.] - ETA: 0s - loss: 0.0155 - acc: 0.9996
8529/8529 [=====] - 79s 9ms/step - loss: 0.0155 - acc: 0.9996 - val_loss: 0.7151 - val_acc: 0.7417
```

Epoch 00020: val_acc did not improve from 0.75949


Evaluating the Model

```
80/2133 [>.....] - ETA: 3s
160/2133 [=>.....] - ETA: 4s
240/2133 [==>.....] - ETA: 3s
320/2133 [===>.....] - ETA: 3s
400/2133 [====>.....] - ETA: 3s
480/2133 [=====>.....] - ETA: 3s
560/2133 [=====>.] - ETA: 3s
640/2133 [=====>.] - ETA: 2s
720/2133 [=====>.] - ETA: 2s
800/2133 [=====>.] - ETA: 2s
880/2133 [=====>.] - ETA: 2s
960/2133 [=====>.] - ETA: 2s
1040/2133 [=====>.] - ETA: 2s
1120/2133 [=====>.] - ETA: 2s
1200/2133 [=====>.] - ETA: 1s
1280/2133 [=====>.] - ETA: 1s
1360/2133 [=====>.] - ETA: 1s
1440/2133 [=====>.] - ETA: 1s
1520/2133 [=====>.] - ETA: 1s
1600/2133 [=====>.] - ETA: 1s
1680/2133 [=====>.] - ETA: 0s
1760/2133 [=====>.] - ETA: 0s
1840/2133 [=====>.] - ETA: 0s
1920/2133 [=====>.] - ETA: 0s
2000/2133 [=====>.] - ETA: 0s
2080/2133 [=====>.] - ETA: 0s
2133/2133 [=====] - 4s 2ms/step
```

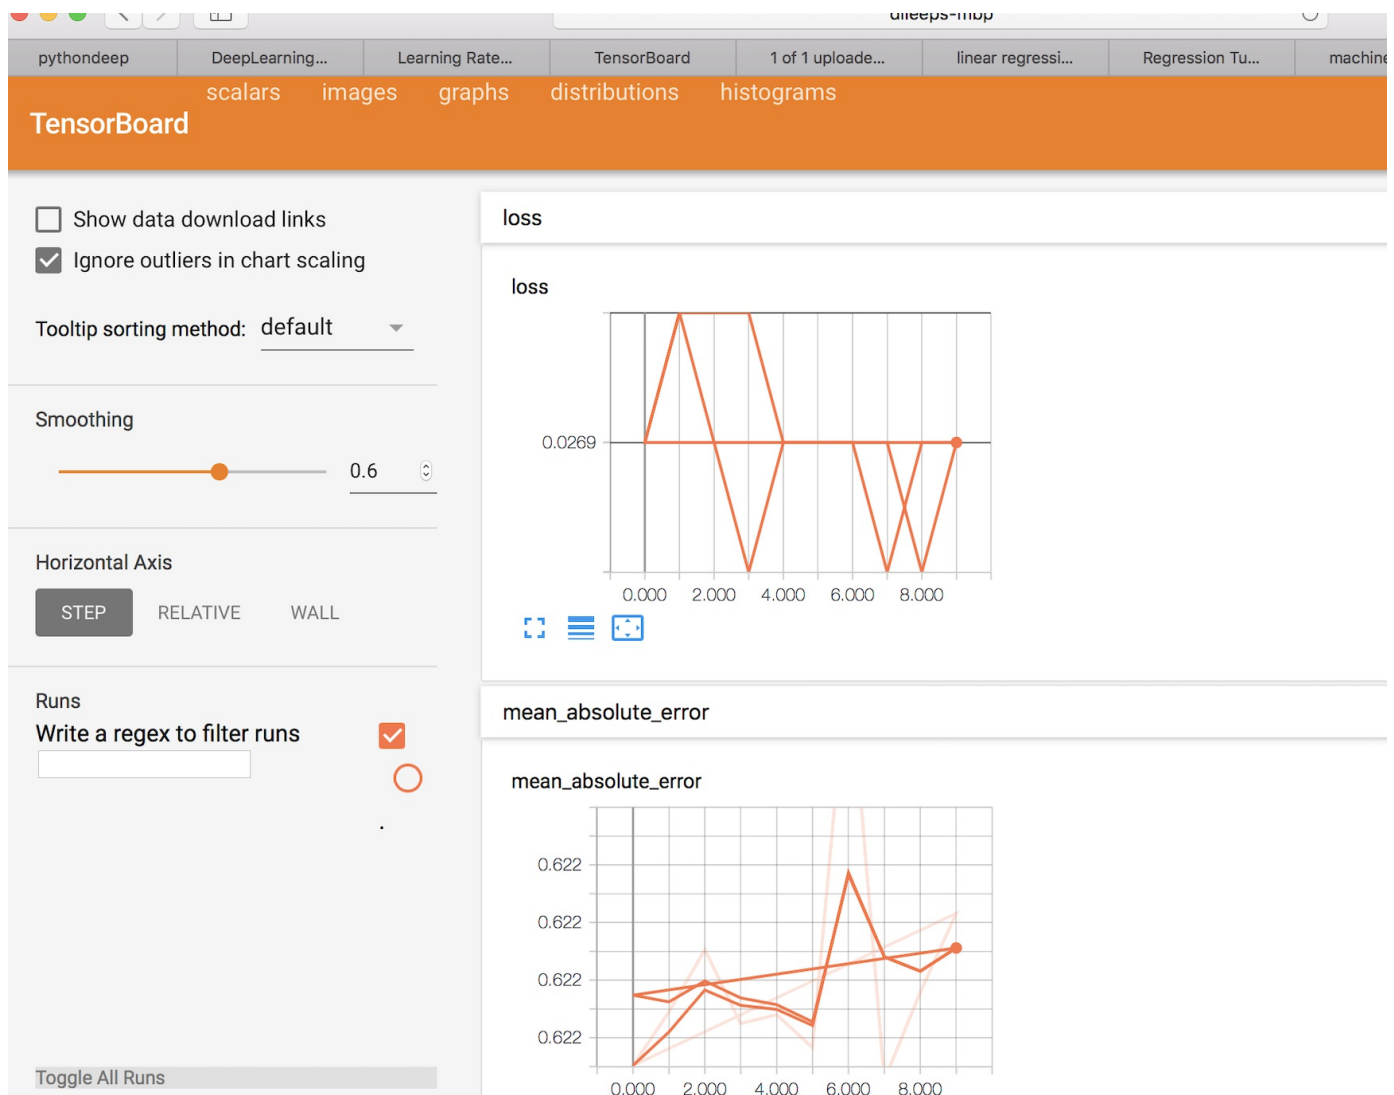
Score: 0.72

Validation Accuracy: 0.74

Figure 1



Tensorboard



5. Text Classification with LSTM

The dataset used is - <https://www.kaggle.com/c/sentiment-analysis-on-movie-reviews/data>

As it is a text data, we used the original dataset and uploaded the content to google colab with runtime as GPU.

The code snippet describing the data is given below

	Phraseld	Sentenceld	Phrase
0	156061	8545	An intermittently pleasing but mostly routine ...
1	156062	8545	An intermittently pleasing but mostly routine ...
2	156063	8545	An
3	156064	8545	intermittently pleasing but mostly routine effort
4	156065	8545	intermittently pleasing but mostly routine


The LSTM model used for Text classification is

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 48, 300)	4120500
lstm_3 (LSTM)	(None, 48, 128)	219648
lstm_4 (LSTM)	(None, 64)	49408
dense_3 (Dense)	(None, 100)	6500
dropout_2 (Dropout)	(None, 100)	0
dense_4 (Dense)	(None, 5)	505
Total params: 4,396,561		
Trainable params: 4,396,561		
Non-trainable params: 0		

We trained the model for 5 epochs, with a batch size of 256 for faster training and got an accuracy of 66%

```
Epoch 4/10
124848/124848 [=====] - 93s 748us/step - loss: 0.6926 - acc: 0.7099 - val_loss: 0.8300 - val_acc: 0.6702
Epoch 5/10
124848/124848 [=====] - 93s 748us/step - loss: 0.6660 - acc: 0.7188 - val_loss: 0.8432 - val_acc: 0.6659
```


The plots showing accuracy for train & validation data is shown below



6. Comparisons on CNN and LSTM models

We can clearly see that by comparing the accuracy and the tensorboard results

CNN text classification has obtained an accuracy of 71.47%

CNN image classification has obtained an accuracy of 31.34%

LSTM has obtained a validation accuracy of 61.29%

CNN text classification though is run on 5, 10, 15 epochs has produced different learning rates with a better accuracy for each epoch with an improving accuracy each time it is feed more data.

LSTM is a slow learner for which it gives varying accuracies for each

time the text is given as input, where it learns for each continuous text data feed.

We can conclude by these accuracies that CNN outperforms LSTM for text classification.